

regression

April 24, 2022

1 Linear Regression

1.0.1 Importing libraries

```
[1]: import pandas as pd
import numpy as np
from datetime import datetime
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score, median_absolute_error
```

1.0.2 Uploading data

```
[2]: url = "https://raw.githubusercontent.com/Agablue-red/Machine-Learning/master/
↳data/Dataset.csv"
data = pd.read_csv(url)
```

```
[3]: data.head()
```

```
[3]:
```

	Date	symbol	sector	score	close	return_rate
0	2004-02-11	SU	Energy Minerals	0.953727	12.830000	NaN
1	2004-02-11	GGG	Producer Manufacturing	0.952753	9.322222	NaN
2	2004-02-11	CWT	Utilities	0.934181	14.245000	NaN
3	2004-02-11	BLL	Process Industries	0.922862	8.012500	NaN
4	2004-02-11	APA	Energy Minerals	0.912117	39.509998	NaN

```
[4]: nan_rows = data[data['return_rate'].isnull()]
if nan_rows.symbol.nunique() == len(nan_rows):
    print("NaN for first period")
```

NaN for first period

Replacing NaNs with 0 value:

```
[5]: data['return_rate'] = data['return_rate'].fillna(0)
```

Looking at the tail of the data, meaning the newest observations:

```
[6]: data.tail()
```

```
[6]:
```

	Date	symbol	sector	score	close	\
30546	2022-02-09	PEP	Consumer Non-Durables	0.701507	171.940002	
30547	2022-02-09	SSNC	Technology Services	0.701123	82.419998	
30548	2022-02-09	GEF	Process Industries	0.697954	56.930000	
30549	2022-02-09	DPZ	Consumer Services	0.697741	444.760010	
30550	2022-02-09	LIFZF	Non-Energy Minerals	0.695644	34.410000	

	return_rate
30546	-0.003189
30547	0.025890
30548	-0.001753
30549	0.015272
30550	0.069630

1.0.3 Information about dataset

Data types:

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30551 entries, 0 to 30550
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date             30551 non-null  object
1   symbol           30551 non-null  object
2   sector           30551 non-null  object
3   score            30551 non-null  float64
4   close            30551 non-null  float64
5   return_rate      30551 non-null  float64
dtypes: float64(3), object(3)
memory usage: 1.4+ MB
```

Checking if there is any lack of data:

```
[8]: data.isnull().sum()
```

```
[8]: Date            0
     symbol          0
     sector          0
     score           0
     close           0
     return_rate     0
     dtype: int64
```

Changing the type of 'date' variable:

```
[9]: data['Date'] = pd.to_datetime(data['Date'], format = '%Y-%m-%d')
data = data.set_index('Date')
```

Fundamental statistics on numeric variables

```
[10]: data.describe()
```

```
[10]:
```

	score	close	return_rate
count	30551.000000	30551.000000	30551.000000
mean	0.731206	101.353658	0.003849
std	0.117692	2627.016498	0.044643
min	0.413554	0.020000	-0.951550
25%	0.653428	26.072500	-0.016298
50%	0.741474	44.770000	0.002865
75%	0.813471	73.910004	0.023672
max	0.987225	453000.000000	0.632911

There are in total 30 551 observations. The mean score for this dataset is 0,73, mean closing price is 101,3 and mean return rate is 0,004.

```
[11]: data.symbol.value_counts()
```

```
[11]: SHW      170
GEF       140
ORLY      138
INGR      122
GPC       122
...
TREV      1
HRNN      1
REGI      1
PRMW      1
DWSN      1
Name: symbol, Length: 1338, dtype: int64
```

There are 1338 companies in total, some of them occur only once in the time series and some even over 100 times.

1.0.4 Splitting the data into training and testing sets

Training set involves data from 2010 to 2020 and testing set includes the year 2021.

```
[12]: X_train = data['2010':'2020'].drop(['symbol', 'sector', 'return_rate', 'close'], axis = 1)
y_train = data.loc['2010':'2020', 'return_rate']

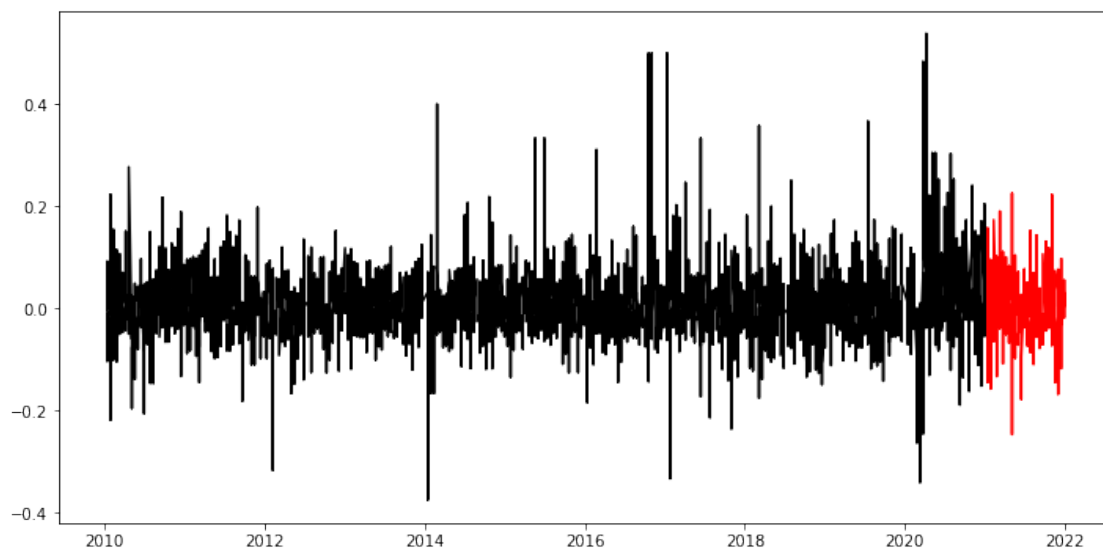
X_test = data['2021'].drop(['symbol', 'sector', 'return_rate', 'close'], axis = 1)
y_test = data.loc['2021', 'return_rate']
```

```
<ipython-input-12-e0a87765a94e>:4: FutureWarning: Indexing a DataFrame with a
datetimelike index using a single string to slice the rows, like
`frame[string]`, is deprecated and will be removed in a future version. Use
`frame.loc[string]` instead.
X_test = data['2021'].drop(['symbol', 'sector', 'return_rate', 'close'], axis =
1)
```

```
[13]: fig, ax=plt.subplots(figsize=(12, 6))

plt.plot(y_train, color = "black")
plt.plot(y_test, color = "red")
```

```
[13]: [<matplotlib.lines.Line2D at 0x24b4770d580>]
```



```
[14]: print("Number transactions X_train dataset: ", X_train.shape)
print("Number transactions y_train dataset: ", y_train.shape)
print("Number transactions X_test dataset: ", X_test.shape)
print("Number transactions y_test dataset: ", y_test.shape)
```

```
Number transactions X_train dataset: (19797, 1)
Number transactions y_train dataset: (19797,)
Number transactions X_test dataset: (2021, 1)
Number transactions y_test dataset: (2021,)
```

1.0.5 Dummy regression

```
[15]: from sklearn.dummy import DummyRegressor
```

```
[32]: # train model
reg_dummy = DummyRegressor(strategy = 'mean').fit(X_train, y_train)

print('Coefficient of determination:', reg_dummy.score(X_train, y_train))
```

Coefficient of determination: 0.0

0% represents a model that does not explain any of the variation in the response variable around its mean.

```
[34]: # predict & evaluate
y_pred_dum = reg_dummy.predict(X_test)

print("Coefficient of determination (R2): %.5f" % r2_score(y_test , y_pred_dum))
print("Mean absolute error (MAE): %.5f" % np.mean(np.absolute(y_pred -
    y_pred_dum)))
print("Residual sum of squares (MSE): %.5f" % mean_squared_error(y_test,
    y_pred_dum))
print("Root mean squared error (RMSE): %.5f" % np.sqrt(metrics.
    mean_squared_error(y_test, y_pred_dum)))
```

Coefficient of determination (R2): -0.00140

Mean absolute error (MAE): 0.00214

Residual sum of squares (MSE): 0.00178

Root mean squared error (RMSE): 0.04214

1.0.6 Linear regression

```
[24]: from sklearn import metrics

# train model
lm = LinearRegression().fit(X_train, y_train)

print('Coefficient of determination:', lm.score(X_train, y_train))
print('Intercept:', lm.intercept_)
print('Slope:', lm.coef_)
```

Coefficient of determination: 0.005437296983874185

Intercept: 0.02420106384509441

Slope: [-0.02744261]

$$f(x) = b x + b$$

$$f(x) = - 0.027x + 0.024$$

$$^2 = 0.0054$$

Model explains only 0.0054 of the variation in the response variable around its mean.

Measure of fit of a model

```
[30]: # predict & evaluate
y_pred = lm.predict(X_test)

print('predicted response:', y_pred, sep='\n')

print("Coefficient of determination (R2): %.5f" % r2_score(y_test , y_pred) )
print("Mean absolute error (MAE): %.5f" % np.mean(np.absolute(y_pred - y_test)))
print("Residual sum of squares (MSE): %.5f" % mean_squared_error(y_test, y_pred))
print("Root mean squared error (RMSE): %.5f" % np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

predicted response:

```
[-0.00099228 -0.00054537 -0.00040295 ...  0.00468725  0.00471852
 0.00475165]
```

Coefficient of determination (R2): -0.00431

Mean absolute error (MAE): 0.03100

Residual sum of squares (MSE): 0.00178

Root mean squared error (RMSE): 0.04220

Adjusted R squared is adjusted for the number of independent variables in the model and equal -0.00431 (adjusted R² will always be less than or equal to R²).

The average of the residuals equal 0.031.

The variance of the residuals equal 0.00178.

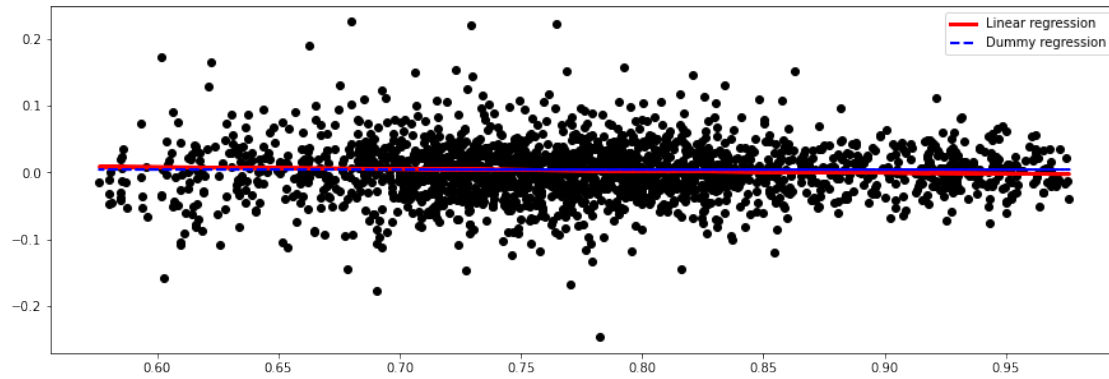
The standard deviation of residuals equal 0.0422.

1.0.7 Comparison between dummy regression and linear regression in combination with observations from test set.

```
[19]: fig, ax=plt.subplots(figsize=(15, 5))

plt.scatter(X_test, y_test, color='black')
plt.plot(X_test, y_pred, color='red', linewidth=3, label='Linear regression')
plt.plot(X_test, y_pred_dum, color='blue', linestyle = 'dashed', linewidth=2,
        label = 'Dummy regression')
ax.legend()
```

```
[19]: <matplotlib.legend.Legend at 0x24b49a5c880>
```



Model does not explain any of the variation in the response variable around its mean.

Linear regression is marginally better than dummy regression.

Both models are not well fit.